

▼ 2a.

Build a Convolutional Neural Network, like what we built in lectures (without skip connections), to classify the images across all 10 classes in CIFAR 10. You need to adjust the fully connected layer at the end properly with respect to the number of output classes. Train your network for 300 epochs. Report your training time, training loss, and evaluation accuracy after 300 epochs. Analyze your results in your report and compare them against a fully connected network (homework 2) on training time, achieved accuracy, and model size.

▼ CIFAR 10 setup:

```
1 import torch
2 import numpy as np
3 from torchvision import datasets
4 import torchvision.transforms as transforms
5 from torch.utils.data.sampler import SubsetRandomSampler
6
7 device = torch.device('cuda:0')
8 print(torch.cuda.is_available())
```

True

```
1 # number of subprocesses to use for data loading
2 num_workers = 0
3
4 # how many samples per batch to load
5 batch_size = 20
6
7 # percentage of training set to use as validation
8 valid_size = 0.2
9
10 # convert data to a normalized torch.FloatTensor
11 transform = transforms.Compose([
12     transforms.ToTensor(),
13     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
14 ])
15
16 # choose the training and test datasets
17 train_data = datasets.CIFAR10('data', train=True,
18                               download=True, transform=transform)
19 test_data = datasets.CIFAR10('data', train=False,
20                               download=True, transform=transform)
21
22 # obtain training indices that will be used for validation
23 num_train = len(train_data)
```

```
24 indices = list(range(num_train))
25 np.random.shuffle(indices)
26 split = int(np.floor(valid_size * num_train))
27 train_idx, valid_idx = indices[split:], indices[:split]
28
29 # define samplers for obtaining training and validation batches
30 train_sampler = SubsetRandomSampler(train_idx)
31 valid_sampler = SubsetRandomSampler(valid_idx)
32
33 # prepare data loaders (combine dataset and sampler)
34 train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
35      sampler=train_sampler, num_workers=num_workers)
36 valid_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
37      sampler=valid_sampler, num_workers=num_workers)
38 test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
39      num_workers=num_workers)
40
41 # specify the image classes
42 classes = ['airplane', 'automobile', 'bird', 'cat', 'deer',
43      'dog', 'frog', 'horse', 'ship', 'truck']
```

Files already downloaded and verified

Files already downloaded and verified

```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3
4 # helper function to un-normalize and display an image
5 def imshow(img):
6     img = img / 2 + 0.5 # unnormalize
7     plt.imshow(np.transpose(img, (1, 2, 0))) # convert from Tensor image
8
9 # obtain one batch of training images
10 dataiter = iter(train_loader)
11 images, labels = dataiter.next()
12 images = images.numpy() # convert images to numpy for display
13
14 # plot the images in the batch, along with the corresponding labels
15 fig = plt.figure(figsize=(25, 4))
16
17 # display 10 images
18 for idx in np.arange(10):
19     ax = fig.add_subplot(2, 10/2, idx+1, xticks=[], yticks=[])
20     imshow(images[idx])
21     ax.set_title(classes[labels[idx]])
```



▼ Model 1 Architecture:

```

1  import torch.nn as nn
2  import torch.nn.functional as F
3
4  # define the CNN architecture
5  class Net1(nn.Module):
6      def __init__(self):
7          super(Net1, self).__init__()
8          self.conv1 = nn.Conv2d(3, 6, 5)
9          self.pool = nn.MaxPool2d(2, 2)
10         self.conv2 = nn.Conv2d(6, 16, 5)
11         self.fc1 = nn.Linear(16 * 5 * 5, 120)
12         self.fc2 = nn.Linear(120, 84)
13         self.fc3 = nn.Linear(84, 10)
14
15     def forward(self, x):
16         x = self.pool(F.relu(self.conv1(x)))
17         x = self.pool(F.relu(self.conv2(x)))
18         x = x.view(-1, 16 * 5 * 5)
19         x = F.relu(self.fc1(x))
20         x = F.relu(self.fc2(x))
21         x = self.fc3(x)
22         return x
23
24     # create a complete CNN
25     model1 = Net1()
26     if torch.cuda.is_available():
27         model1.cuda()
28     print(model1)
29

```

```

Net1(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)

```

```
1 import torch.optim as optim
2
3 # specify loss function
4 criterion = nn.CrossEntropyLoss()
5
6 # specify optimizer
7 optimizer = optim.SGD(model1.parameters(), lr=.01)
```

▼ Model 1: training & testing

```
1 %%time
2 from tqdm import tqdm
3
4 # number of epochs to train the model
5 n_epochs = 300
6
7 #List to store loss to visualize
8 train_losslist = []
9
10 # track change in validation loss
11 valid_loss_min = np.Inf
12
13 for epoch in range(1, n_epochs+1):
14     print('\nEpoch: {}'.format(epoch))
15     # keep track of training and validation loss
16     train_loss = 0.0
17     valid_loss = 0.0
18
19     print('training: ')
20     with tqdm(train_loader, unit="batch") as tepoch:
21         # train the model
22         model1.train()
23         for data, target in tepoch:
24
25             # move tensors to GPU if CUDA is available
26             data, target = data.to(device), target.to(device)
27
28             # clear the gradients of all optimized variables
29             optimizer.zero_grad()
30
31             # forward pass: compute predicted outputs by passing inputs to the model
32             output = model1(data)
33
34             # calculate the batch loss
35             loss = criterion(output, target)
36
37             # backward pass: compute gradient of the loss with respect to model parameters
38             loss.backward()
39
```

```

40         # perform a single optimization step (parameter update)
41         optimizer.step()
42
43         # update training loss
44         train_loss += loss.item()*data.size(0)
45
46     print('validation: ')
47     with tqdm(valid_loader, unit="batch") as vepoch:
48         # validate the model
49         model1.eval()
50         for data, target in vepoch:
51
52             # move tensors to GPU if CUDA is available
53             data, target = data.to(device), target.to(device)
54
55             # forward pass: compute predicted outputs by passing inputs to the model
56             output = model1(data)
57
58             # calculate the batch loss
59             loss = criterion(output, target)
60
61             # update average validation loss
62             valid_loss += loss.item()*data.size(0)
63
64     # calculate average losses
65     train_loss = train_loss/len(train_loader.dataset)
66     valid_loss = valid_loss/len(valid_loader.dataset)
67     train_losslist.append(train_loss)
68
69     # print training/validation statistics
70     print('Training Loss: {:.6f} \tValidation Loss: {:.6f}'.format(train_loss, valid_loss))
71
72     # save model if validation loss has decreased
73     if valid_loss <= valid_loss_min:
74         print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
75             valid_loss_min,
76             valid_loss))
77         torch.save(model1.state_dict(), 'model1_cifar.pt')
78         valid_loss_min = valid_loss
79
80 # plt.plot(n_epochs, train_losslist)
81 # plt.xlabel("Epoch")
82 # plt.ylabel("Loss")
83 # plt.title("Performance of Model 1")
84 # plt.show()

```

```

Epoch: 1
training:
100%|██████████| 2000/2000 [00:11<00:00, 181.50batch/s]
validation:
100%|██████████| 500/500 [00:02<00:00, 223.27batch/s]

```

Training Loss: 1.757462 Validation Loss: 0.389919
 Validation loss decreased (inf --> 0.389919). Saving model ...

Epoch: 2
 training:
 100%|██████████| 2000/2000 [00:11<00:00, 176.81batch/s]
 validation:
 100%|██████████| 500/500 [00:02<00:00, 231.08batch/s]
 Training Loss: 1.421384 Validation Loss: 0.321831
 Validation loss decreased (0.389919 --> 0.321831). Saving model ...

Epoch: 3
 training:
 100%|██████████| 2000/2000 [00:10<00:00, 182.56batch/s]
 validation:
 100%|██████████| 500/500 [00:02<00:00, 215.26batch/s]
 Training Loss: 1.234788 Validation Loss: 0.288718
 Validation loss decreased (0.321831 --> 0.288718). Saving model ...

Epoch: 4
 training:
 100%|██████████| 2000/2000 [00:10<00:00, 181.92batch/s]
 validation:
 100%|██████████| 500/500 [00:02<00:00, 230.61batch/s]
 Training Loss: 1.134695 Validation Loss: 0.270984
 Validation loss decreased (0.288718 --> 0.270984). Saving model ...

Epoch: 5
 training:
 100%|██████████| 2000/2000 [00:11<00:00, 181.71batch/s]
 validation:
 100%|██████████| 500/500 [00:02<00:00, 234.99batch/s]
 Training Loss: 1.058859 Validation Loss: 0.255602
 Validation loss decreased (0.270984 --> 0.255602). Saving model ...

Epoch: 6
 training:
 100%|██████████| 2000/2000 [00:10<00:00, 184.02batch/s]
 validation:
 100%|██████████| 500/500 [00:02<00:00, 230.67batch/s]
 Training Loss: 0.993789 Validation Loss: 0.243864
 Validation loss decreased (0.255602 --> 0.243864). Saving model ...

Epoch: 7
 training:
 100%|██████████| 2000/2000 [00:10<00:00, 182.31batch/s]
 validation:
 100%|██████████| 500/500 [00:02<00:00, 231.69batch/s]
 Training Loss: 0.941940 Validation Loss: 0.231812
 Validation loss decreased (0.243864 --> 0.231812). Saving model ...

```
1 model1.load_state_dict(torch.load('model1_cifar.pt'))
```

<All keys matched successfully>

```

1 # track test loss
2 test_loss = 0.0
3 class_correct = list(0. for i in range(10))
4 class_total = list(0. for i in range(10))
5
6 model1.eval()
7 # iterate over test data
8 for data, target in test_loader:
9     # move tensors to GPU if CUDA is available
10    data, target = data.to(device), target.to(device)
11    # forward pass: compute predicted outputs by passing inputs to the model
12    output = model1(data)
13    # calculate the batch loss
14    loss = criterion(output, target)
15    # update test loss
16    test_loss += loss.item()*data.size(0)
17    # convert output probabilities to predicted class
18    _, pred = torch.max(output, 1)
19    # compare predictions to true label
20    correct_tensor = pred.eq(target.data.view_as(pred))
21    correct = np.squeeze(correct_tensor.cpu().numpy())
22    # calculate test accuracy for each object class
23    for i in range(batch_size):
24        label = target.data[i]
25        class_correct[label] += correct[i].item()
26        class_total[label] += 1
27
28 # average test loss
29 test_loss = test_loss/len(test_loader.dataset)
30 print('Test Loss: {:.6f}\n'.format(test_loss))
31
32 for i in range(10):
33     if class_total[i] > 0:
34         print('Test Accuracy of %5s: %2d%% (%2d/%2d)' % (
35             classes[i], 100 * class_correct[i] / class_total[i],
36             np.sum(class_correct[i]), np.sum(class_total[i])))
37     else:
38         print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))
39
40 print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
41     100. * np.sum(class_correct) / np.sum(class_total),
42     np.sum(class_correct), np.sum(class_total)))

```

Test Loss: 1.030446

Test Accuracy of airplane: 73% (732/1000)
 Test Accuracy of automobile: 75% (752/1000)
 Test Accuracy of bird: 45% (457/1000)
 Test Accuracy of cat: 40% (402/1000)
 Test Accuracy of deer: 60% (601/1000)

```
Test Accuracy of  dog: 67% (670/1000)
Test Accuracy of  frog: 65% (656/1000)
Test Accuracy of horse: 75% (754/1000)
Test Accuracy of  ship: 76% (767/1000)
Test Accuracy of truck: 71% (712/1000)
```

```
Test Accuracy (Overall): 65% (6503/10000)
```

▼ 2b.

Extend your CNN by adding one more additional convolution layer followed by an activation function and pooling function. You also need to adjust your fully connected layer properly with respect to intermediate feature dimensions. Train your network for 300 epochs. Report your training time, loss, and evaluation accuracy after 300 epochs. Analyze your results in your report and compare your model size and accuracy over the baseline implementation in Problem1.a. Do you see any over-fitting? Make sure to submit your code by providing the GitHub URL of your course repository for this course.

► Model 2 Architecture:

[] ↳ 2 cells hidden

► Model 2: training & testing

[] ↳ 3 cells hidden

